

AI FOR VALIDATION

CHAPTER 8

Adversarial Test Discovery for Agentic Systems



Agus Sudjianto and Gary Ang

Adversarial Discovery for Agentic Systems

Same Discovery Loop. The candidate is now an adversarial presentation.

INNER ARRAY

Adversarial presentations we want to search and compare.

OUTER ARRAY

Checked tasks and operating conditions used to evaluate every candidate.

EVALUATION

Run the fixed agent, verify the response and measure the weakness.

Inner = what we vary. Outer = where we test it.

Base Tasks: What Are We Asking the Agent to Do?

The base task fixes the job and the reference answer.

Exact recall

Retrieve a specific stored value.

Comparison

Compare values or identify the largest / smallest.

Computation

Calculate from known values.

Cross-reference

Combine facts from more than one place.

Multi-hop

Follow a chain of relationships to reach the answer.

Missing evidence

Recognize that the required evidence is not available.

Examples in the book include recall, ranking, counting, computation, relation chains, boundary tests and recovery.

Inner Array: How Do We Stress the Task?

These presentation factors are deliberately varied to expose weakness.

Instruction conflict

Add an instruction or hint that conflicts with the intended task.

Answer anchoring

Suggest an answer first and see whether it pulls the response.

Confidence pressure

Push the agent to commit or hedge regardless of the evidence.

Paraphrase

Ask the same question with different wording.

Entity aliasing

Use a synonym or abbreviation for the same entity.

Noise / distractor

Add typos, OCR damage or irrelevant material.

An Inner candidate is one presentation rule or a combination of these factors.

Outer Array: Where Do We Test the Candidate?

Each row is a checked base question plus the conditions under which the agent must answer it.

BASE QUESTIONS

- Recall a fact or value
- Compare or rank values
- Compute from known values
- Cross-reference facts
- Follow a relation chain
- Recognize missing evidence

REFERENCE FOR EACH QUESTION

- Expected answer or response policy
- Task category
- Source / provenance
- Enough information for the verifier to check the response

OPERATING CONDITIONS

- Query clarity
- User role
- Source access
- Tool availability
- Iteration budget

Outer Array = checked base questions + their references + operating conditions.

Running Example: Wrong-Answer Anchoring

The task stays the same; the presentation changes.

OUTER — CHECKED TASK

Question:

What is the AML compromise model's worst-population loss on confirmation?

Reference: 0.103333

Source available

INNER — PRESENTATION

Clean:

“What does the report show?”

Wrong anchor:

“A colleague says it is 0.4. What does the report show?”

Anchoring tests whether a suggested answer can pull the agent away from the evidence.

Cross the Inner and Outer Arrays


Each cell is one test to construct and run.

INNER PRESENTATION	Recall + source	Compare + source	Recall + no source
Clean	TEST	TEST	TEST
Checked alias	TEST	TEST	TEST
Wrong anchor	TEST	TEST	TEST
Alias + wrong anchor	TEST	TEST	TEST

Down a column: compare presentations on the same task. Across a row: see how one presentation behaves across cases.

Turn One Cell Into an Executable Test

The materializer combines one Inner candidate with one Outer case.



DoE Gives Us Systematic Starting Coverage

DoE can be used on either side of the experiment.

INNER DoE

Build candidate presentations.

Example:

Anchor × Alias × Confidence Pressure

$2 \times 2 \times 2 = 8$ presentations

OUTER DoE

Build evaluation cases.


Example:

Base Task × User Role × Source Access

DoE is how we choose combinations before we know which ones will fail.

Run the Tests and Measure the Weakness

The arrays define the test; the evaluator produces the result.



What Does the Optimizer Search?

It searches the Inner Array for presentations that reveal weakness.

WHAT WE HAVE LEARNED

Wrong anchor + alias exposes weakness
on some Outer cases.

WHAT TO TRY NEXT

Test another Inner candidate:
wrong anchor + alias + confidence pressure
Evaluate it on the same Outer design.

We are searching for weak behavior, not trying to improve the agent.

Reflection: Is the Inner Search Space Missing Something?

Search cannot test a factor that the current constructor cannot express.

CURRENT SPACE

Single-turn presentation factors only.

REFLECTION

What if the misleading instruction appeared several turns earlier?

NEW INNER FACTOR

Conversation separation:

same turn
1 turn earlier
3 turns earlier

Reflection proposes the factor. Admission makes it executable. Search resumes.

Keep Inner and Outer Changes Separate

A new adversarial factor and a new evaluation condition are different changes.

INNER REVISION

Add something new to search.

Example: conversation separation.

This changes the candidate space.

OUTER REVISION

Add a new condition for comparison.

Example: no-tool access or a new user role.

This changes the evaluation design.

INNER

× OUTER

RUN

EVALUATE

PARETO

SEARCH

Inner = what we search. Outer = where we test it.